

Contact

91 9769545966 (Mobile)
vijaykumar.tater@gmail.com

www.linkedin.com/in/vijaykumarjain
(LinkedIn)

Top Skills

PostgreSQL
Puppet
python

Languages

English (elementary)
Hindi (elementary)

vijaykumar jain

Looking for Infrastructure Engineering related openings.
Mumbai, Maharashtra, India

Summary

Software Engineer with strong knowledge of application/system software/infrastructure monitoring principles and practice along with reading writing code to enable products release features faster.

Experience

OpenTable

Database SRE

June 2017 - November 2019 (2 years 6 months)

Mumbai Area, India

Postgresql (11, 10, 9.6), pgbouncer, haproxy, sensu, graphite, filebeat, ELK, python, psycpg2

Providing day to day support to more than 100 postgres clusters.

Setup Postgres cluster (one primary and multiple hot standby for reads) using puppet on vmware.

Make use of PGBouncer connection pooling.

Scale the read nodes by putting them behind HAProxy with health check as replication lag.

Failover or promote a hot standby to primary.

Taking database dumps and restoring them.

Performed migration of more than 15 clusters from 9.6 to 10, with both clusters active (with minimum downtime for writes) in the event of rollback.

Sending postgresql statistics metrics via collectd to graphite to understand database bottlenecks and growth patterns.

Setting up monitoring and alerting for postgresql via sensu.

Sending postgresql logs via filebeat to logstash to be viewed on Kibana.

Using Barman to configure point in time recovery for databases.

Puppetizing the sharded postgres cluster (3.5TB) using foreign data wrappers and puppet exported resources.

Being able to take over primary support for postgres for SF team when the core postgres DBA was on leave.

Understand the multimaster mesh structure across zones using pg10 logical replication.

Using logical replication mechanism to upgrade a postgres cluster.

Use of explain to understand query bottlenecks.

Setup postgres 11 POC cluster, to work on partitioning for historical databases.

Tuning the database (for performance and high upserts via vacuum) based on pgtune and autovacuum.

Use of hot_standby_feedback to ensure long running queries complete on standby without autovacuum cancelling them.

Attended PGCon2018 and follow blogs.

Mongo:

Basic monitoring of mongo clusters using cloud mms monitoring.

Good awareness of mongo replication and using mongo sharding/mongos.

mongodump and mongo restore of collections and databases.

Good awareness of indexing and using explain to figure out long running queries making use of bad plans.

OpenTable

SRE

June 2016 - November 2019 (3 years 6 months)

Mumbai Area, India

Tools: Puppet, Sensu, Graphite, Stasd, ELK, Teamcity, Jenkins, Docker, Pagerduty, Hipchat, Jira, Mesosphere, Singularity, Service Discovery, Kafka, Redis, RabbitMQ, Zookeeper, Postgresql, Mongo, Virtualbox, Vagrant.

Languages: ruby, python, perl and occasionally reading stackoverflow for nodejs/go/powershell.

Automation Expert.

Contributed to opensource sensu and sensu-puppet, sensu-filebeat, and a lot of In-House Puppet Modules.

Using foreman (GUI and API) to provision hosts.

Sending sensu events to logstash using filebeat to view event history on Kibana.

Full ownership of sensu for monitoring OT infrastructure (housing more than 2K servers ubuntu and windows)

POC for migration of PagerDuty to VictorOps.

Diving in source code of tools such as Sensu/Graphite/Postgresql/Mesos to figure out how things work when we have a non reported issue.

Awareness of microservices architecture and work in a team to support Mesos environment with custom Service Discovery and Singularity as Task Scheduler.

The desire to explore open source tooling with scalable architecture with apt monitoring and alerting infrastructure to make engineering life better, I say it loud that I am proud of being Ops Support.

Nomura

Senior Infrastructure Engineer

September 2011 - May 2016 (4 years 9 months)

Mumbai

Development Lead of the UnixEngineering Team.

Provided a comprehensive JIRA workflow for the team to use for puppet development.

Test Driven Development using puppet, rspec, serverspec, beaker, rvm, git.

Responsible for POC implementation of GitLab for SCM, and successfully helped to move it to production.

Setup Continuous Integration for Redhat Satellite 6 using Jenkins.

Provided customized CLI in python to use pulp (OSS) to replace Opsware for RPM management and deployment.

Create RPMs using fpm tool.

Automated existing OS build process using autotest.

Worked on creating a REST-ful API for LDAP data management.

Provided Training in the firm on python, git, and basic unix

Setup SSL for web based applications like GitLab and pulp in the firm.

Took active interest in sysadmin activities like provisioning and building

Operating Systems, Network Boot from

consoles, KVM, Vagrant, VSphere, Kerberos.

Basic knowledge of SQL and NoSQL databases.

Basic knowledge of messaging services like qpid.

Working knowledge of Moose and DBIx::Class for development.

Contributed to open source projects such as

Pulp (<https://github.com/pulp/pulp>)

autotest (<https://github.com/autotest/autotest>)

Multiprocessing/Threading(No shared states)

Exposure of using LDAP/Sybase databases using perl.

Using GitHub, GitLab.

Attended Redhat, Vmware, Amazon events.

Self training on modern technologies using PluralSight and Youtube.

Bank of America Merrill Lynch

Senior Software Engineer

November 2009 - June 2012 (2 years 8 months)

Analysis of enhancement requests passed on to the offshore team.

Use Perl Best Practices (PBP) to ensure code maintainability and readability.

Implementation of Moose and Moosex Modules in Perl to support PBP.

Perform peer review of code written by other team members.

Setting up of Geneos (monitoring tool) to reduce monitoring for support team

Provided a solution to maintain best archive for tick files collected on multiple regions and thereby helping in removing redundancy.

Analyse system crashes from core dumps using strace and truss/pstack/gdb.

Automated batch monitoring to reduce manual intervention for L1 Support via ITRS.

Installation and use of mon (MDA) tables in Sybase to aid monitoring.

Real time validation for Vhayu and XMIM tick level data and immediate alerting on level breach.

Extensive use of ITRS monitoring tool to automate batch, services, process and

Feed files collection on windows and Linux hosts for monitoring, and implementing corrective actions on the alerts to nullify any manual intervention.

Oracle Financial Services Software Ltd

3 years 1 month

Associate Consultant

August 2007 - November 2009 (2 years 4 months)

Mumbai

Analysis of enhancement requests passed on to the offshore team via SourceForge, JIRA

Development of the requested features using perl and shell scripting with Sybase Database.

Testing of code written by peers

Perform peer review of code written by other team members.

Development of new Stored Procedures, views as per the specifications on Sybase 12.5

Performance tuning of the existing Stored Procedures/sqls using showplan, statistics io , dbcc traceon

(3604,302,310,317)

Writing perl/ksh programs for Batch applications with isql.

Develop perl/korn shell scripts to execute the batch applications.

Use of perl to start jobs in forked mode and monitor load and resource utilization using uptime.

Successfully modifying scripts which used to fail with out of memory errors in perl to write data structures to disk and use less memory.

Auto restart of failed jobs to minimize manual delays.

Use of abstract plans to analyze sql performances.

Associate consultant

November 2006 - July 2007 (9 months)

production support.

Education

Thakur College Of Engineering and Technology

Rajhans Vidyalaya